

DSC 680 Project 3 Red Wine Python

May 10, 2020

```
[16]: import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
# from sklearn import preprocessing
# from sklearn.ensemble import RandomForestRegressor
# from sklearn.pipeline import make_pipeline
# from sklearn.model_selection import GridSearchCV
from sklearn.metrics import mean_squared_error, r2_score
from sklearn.linear_model import LinearRegression
from sklearn import metrics
import seaborn as sns
import matplotlib.pyplot as plt
```

```
[2]: # Read the data
red = pd.read_csv('winequality-red.csv', sep = ';')
red.head()
```

```
[2]:    fixed acidity  volatile acidity  citric acid  residual sugar  chlorides \
0             7.4              0.70         0.00             1.9       0.076
1             7.8              0.88         0.00             2.6       0.098
2             7.8              0.76         0.04             2.3       0.092
3            11.2              0.28         0.56             1.9       0.075
4             7.4              0.70         0.00             1.9       0.076
```

```
    free sulfur dioxide  total sulfur dioxide  density  pH  sulphates \
0              11.0              34.0  0.9978  3.51       0.56
1              25.0              67.0  0.9968  3.20       0.68
2              15.0              54.0  0.9970  3.26       0.65
3              17.0              60.0  0.9980  3.16       0.58
4              11.0              34.0  0.9978  3.51       0.56
```

```
    alcohol  quality
0       9.4        5
1       9.8        5
2       9.8        5
3       9.8        6
4       9.4        5
```

```
[3]: print(red.shape)
```

```
(1599, 12)
```

```
[4]: red.describe()
```

```
[4]:
```

	fixed acidity	volatile acidity	citric acid	residual sugar	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	8.319637	0.527821	0.270976	2.538806	
std	1.741096	0.179060	0.194801	1.409928	
min	4.600000	0.120000	0.000000	0.900000	
25%	7.100000	0.390000	0.090000	1.900000	
50%	7.900000	0.520000	0.260000	2.200000	
75%	9.200000	0.640000	0.420000	2.600000	
max	15.900000	1.580000	1.000000	15.500000	

	chlorides	free sulfur dioxide	total sulfur dioxide	density	\
count	1599.000000	1599.000000	1599.000000	1599.000000	
mean	0.087467	15.874922	46.467792	0.996747	
std	0.047065	10.460157	32.895324	0.001887	
min	0.012000	1.000000	6.000000	0.990070	
25%	0.070000	7.000000	22.000000	0.995600	
50%	0.079000	14.000000	38.000000	0.996750	
75%	0.090000	21.000000	62.000000	0.997835	
max	0.611000	72.000000	289.000000	1.003690	

	pH	sulphates	alcohol	quality
count	1599.000000	1599.000000	1599.000000	1599.000000
mean	3.311113	0.658149	10.422983	5.636023
std	0.154386	0.169507	1.065668	0.807569
min	2.740000	0.330000	8.400000	3.000000
25%	3.210000	0.550000	9.500000	5.000000
50%	3.310000	0.620000	10.200000	6.000000
75%	3.400000	0.730000	11.100000	6.000000
max	4.010000	2.000000	14.900000	8.000000

```
[5]: # Search for missing values
print(red.isnull().sum())
```

```
fixed acidity      0
volatile acidity   0
citric acid        0
residual sugar     0
chlorides          0
free sulfur dioxide 0
total sulfur dioxide 0
density            0
pH                 0
```

```
sulphates          0
alcohol            0
quality            0
dtype: int64
```


```
[6]: # Get column names
col_red_names = red.columns
print(col_red_names)
# Get column data types
print(red.dtypes)
```

```
Index(['fixed acidity', 'volatile acidity', 'citric acid', 'residual sugar',
       'chlorides', 'free sulfur dioxide', 'total sulfur dioxide', 'density',
       'pH', 'sulphates', 'alcohol', 'quality'],
      dtype='object')
fixed acidity      float64
volatile acidity   float64
citric acid        float64
residual sugar     float64
chlorides          float64
free sulfur dioxide float64
total sulfur dioxide float64
density            float64
pH                 float64
sulphates          float64
alcohol            float64
quality            int64
dtype: object
```

```
[7]: correlations = red.corr()['quality'].drop('quality')
print(correlations)
```

```
fixed acidity      0.124052
volatile acidity   -0.390558
citric acid        0.226373
residual sugar     0.013732
chlorides          -0.128907
free sulfur dioxide -0.050656
total sulfur dioxide -0.185100
density            -0.174919
pH                 -0.057731
sulphates          0.251397
alcohol            0.476166
Name: quality, dtype: float64
```

```
[8]: sns.heatmap(red.corr())
plt.show()
```



```
[11]: # Separate target from training features
```

```
y = red.quality
x = red.drop('quality', axis=1)
```

```
[12]: # Split data into train and test sets
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,random_state=3)
```

```
[13]: # fitting linear regression to training data
```

```
regressor = LinearRegression()
regressor.fit(x_train,y_train)
```

```
# this gives the coefficients of the 10 features selected above.
print(regressor.coef_)
```

```
[ 2.32736478e-02 -9.91535569e-01 -1.41267594e-01  8.11925585e-03
 -1.59192407e+00  5.50005690e-03 -3.54198549e-03 -6.06916616e+00
 -4.06325022e-01  8.23603060e-01  2.94180891e-01]
```

```
[14]: # To predict the quality of wine with this model
```

```
train_pred = regressor.predict(x_train)
print(train_pred)
test_pred = regressor.predict(x_test)
print(test_pred)
```

```
[5.33390209 5.33458216 5.94987004 ... 6.39109929 6.20184044 5.27719203]
[5.09908272 5.65580865 5.90927233 6.13810421 5.00495043 5.44066916
 5.05213654 6.15418124 5.52055599 5.77519663 5.61796132 5.23498287
 5.23127869 5.31466808 6.46439345 5.04000017 5.85280918 5.19300859
 6.0919118 6.34255254 6.41600994 5.52588684 5.80534686 4.93255733
 5.16159004 5.48207651 5.13834113 6.59480979 5.89478275 5.73709
 6.09133736 6.29529369 4.91616391 5.88376873 5.10515437 5.96400538
 6.80732578 5.03724291 5.25485064 5.88376873 5.17431803 4.84899008
 6.4903037 5.40465942 5.30375415 5.83513199 5.70825368 5.23988973
 5.24870634 5.46267267 5.08516492 5.61701512 6.01804854 6.32751521
 5.4628648 5.36127481 5.10151339 4.92009423 5.2240759 5.08722001
 4.79258875 5.43567381 5.25054561 5.6798788 5.85050157 6.52603804
 5.37941315 5.71598525 5.16966353 5.98159839 5.63912543 5.6004759
 5.74068429 5.22739422 5.98184324 5.51332746 5.40647057 5.68342011
 5.64578506 5.73709 6.23278066 5.29710528 4.66398697 6.0425789
 5.53767287 5.17796008 5.21203744 5.95953904 5.51273214 5.64429718
 5.70470381 5.64311292 5.72484629 5.31747436 5.37088603 5.40115158
 4.81676854 5.44991141 5.47380406 6.536759 6.14384914 5.63963684
 6.0764213 6.18184545 5.7333969 4.93007305 4.7323061 5.04093649
 5.45014754 5.78173092 6.44429962 5.47723449 6.46843631 5.94017696
 5.43087432 5.2047468 5.3484345 5.20416986 6.19702214 5.62516963
 5.83959227 5.19657408 5.17215817 5.26956927 5.74449907 5.64311292
 6.14688619 5.89603605 5.49258006 5.40446964 5.25143958 5.40731851
 5.70426783 5.67629516 6.58239965 5.88935979 6.38222076 5.73042293
 5.37034033 5.14004751 5.58608223 6.59675055 5.24400122 5.25667208
 5.54847605 5.16833418 5.76174468 6.10180871 6.93344996 4.99041658
 5.02045525 4.6936361 5.83279744 5.0202852 5.22739422 5.70787157
 5.629781 5.33023534 5.22441656 5.84507489 5.6222602 5.78080916
 5.52275613 6.07368903 5.63220544 5.49065808 5.96678121 4.82131562
 5.25184801 5.67298643 5.7315135 6.61949155 5.0208738 5.90323477
 5.85139093 5.20850319 5.68647856 5.5107533 5.40465942 6.37989991
 6.72172128 4.98279688 5.87260366 5.75498799 5.74800561 5.61884195
 5.711974 5.42385235 6.05973318 5.58256745 5.8850805 6.52201233
 5.00475912 5.39601068 5.1816475 5.16966353 5.4628648 5.75133323
 5.69516241 4.92576461 5.12771702 4.98053813 6.18910436 5.66320044
 5.44465486 5.56328176 4.99497389 5.83781624 5.31649795 5.48545825
 5.69404608 5.6982245 5.70724638 5.82340331 5.79368731 6.02336481
 6.20271572 5.27399606 5.04145317 5.21266809 5.38322644 4.97897504
 6.20624438 5.44065494 5.94166869 5.2181219 6.61288079 5.08900422
 5.29347654 5.03590649 6.17278277 5.77436736 4.8608174 5.72198789]
```

```

5.29496071 5.35533595 5.17179335 6.29984401 5.59779315 4.9609217
6.10214544 6.03200835 6.17804629 5.42115437 6.76301676 6.21226379
6.0861762 5.22659704 5.44997449 5.54506957 5.35533595 5.2412872
5.74150467 5.25027186 6.12552958 5.42771411 5.83939322 4.81992631
6.06232111 5.08298843 6.43084951 6.06506094 5.70218957 5.70426783
4.90078952 5.99868147 5.28886652 5.70122979 5.42595998 5.11848917
6.48395933 5.30489074 5.96848836 5.67858152 6.4906849 6.19855569
5.09593756 5.46687509 5.30216736 5.23795665 6.38222076 5.37816899
5.40613284 5.99854543 4.99943884 5.88539114 5.3527035 5.18049126
5.48464387 5.91747035 4.82427368 6.84444638 5.17024164 4.93584868
5.71503791 5.67672479 5.29391086 6.28837157 6.88282676 6.58112403
5.94280234 6.33341072 5.90237267 5.56893141 6.00359575 5.51815475
5.54851918 5.7333969 5.31204514 5.1501443 6.2167232 5.30100053
6.21938621 6.10351244 5.87211737 5.4294743 4.84344446 6.0468106
5.24475561 5.11875658 6.4639068 5.49851347 6.74724534 5.10901438
5.1652877 5.30672745 5.71398551 5.10794561 5.53117148 5.31620746
5.16515101 4.97232263 5.45290614 5.38377742 5.3484098 5.17608733
6.11753324 5.61279425 5.73151521 6.34758037 5.90725513 5.5107533
5.69333789 5.14471327 5.72896485 6.40189738 6.15898319 5.46033547
6.05278052 5.73135561 6.23901028 5.23795665 5.54774909 5.04649062
4.9961328 5.07773137 5.80157819 5.39684518 6.01142215 5.14597383
6.51552207 5.3931712 5.49159322 5.64862532 5.52478877 5.26990054
6.4906849 5.82378383 5.03571259 5.24297506 5.49335 5.28186183
5.50973691 5.04135168 5.29294649 5.45465132 5.36525182 6.12969147
4.99593587 5.30216736 6.07959665 5.2047468 5.13404621 4.65965123
6.15655748 6.15713628 6.50577148 5.79958773 5.72198789 6.39773211
6.13733354 5.89478275 6.05255784 6.03220643 5.37428389 5.40499404
5.6190012 5.40207228 5.75989893 5.25849574]

```

```

[17]: # calculating rmse
train_rmse = mean_squared_error(train_pred, y_train) ** 0.5
print(train_rmse)
test_rmse = mean_squared_error(test_pred, y_test) ** 0.5
print(test_rmse)

```

```

0.6524682504422629
0.6269476348621655

```

```

[18]: # rounding off the predicted values for test set
predicted_data = np.round_(test_pred)
print(predicted_data)
print('Mean Absolute Error:', metrics.mean_absolute_error(y_test, test_pred))
print('Mean Squared Error:', metrics.mean_squared_error(y_test, test_pred))
print('Root Mean Squared Error:', np.sqrt(metrics.mean_squared_error(y_test,
↪test_pred)))

```

```

[5. 6. 6. 6. 5. 5. 5. 6. 6. 6. 6. 5. 5. 5. 6. 5. 6. 5. 6. 6. 6. 6. 5.
 5. 5. 5. 7. 6. 6. 6. 6. 5. 6. 5. 6. 7. 5. 5. 6. 5. 5. 6. 5. 5. 6. 6. 5.]

```

```

5. 5. 5. 6. 6. 6. 5. 5. 5. 5. 5. 5. 5. 5. 5. 6. 6. 7. 5. 6. 5. 6. 6. 6.
6. 5. 6. 6. 5. 6. 6. 6. 6. 5. 5. 6. 6. 5. 5. 6. 6. 6. 6. 6. 5. 5. 5.
5. 5. 5. 7. 6. 6. 6. 6. 6. 5. 5. 5. 5. 6. 6. 5. 6. 6. 5. 5. 5. 5. 6. 6.
6. 5. 5. 5. 6. 6. 6. 6. 5. 5. 5. 5. 6. 6. 7. 6. 6. 6. 5. 5. 6. 7. 5. 5.
6. 5. 6. 6. 7. 5. 5. 5. 6. 5. 5. 6. 6. 5. 5. 6. 6. 6. 6. 6. 5. 6. 5.
5. 6. 6. 7. 5. 6. 6. 5. 6. 6. 5. 6. 7. 5. 6. 6. 6. 6. 6. 5. 6. 6. 6. 7.
5. 5. 5. 5. 5. 6. 6. 5. 5. 5. 6. 6. 5. 6. 5. 6. 5. 5. 6. 6. 6. 6. 6. 6.
6. 5. 5. 5. 5. 5. 6. 5. 6. 5. 7. 5. 5. 5. 6. 6. 5. 6. 5. 5. 5. 6. 6. 5.
6. 6. 6. 5. 7. 6. 6. 5. 5. 6. 5. 5. 6. 5. 6. 5. 6. 5. 6. 5. 6. 6. 6. 6.
5. 6. 5. 6. 5. 5. 6. 5. 6. 6. 6. 6. 5. 5. 5. 5. 6. 5. 5. 6. 5. 6. 5. 5.
5. 6. 5. 7. 5. 5. 6. 6. 5. 6. 7. 7. 6. 6. 6. 6. 6. 6. 6. 5. 5. 6. 5.
6. 6. 6. 5. 5. 6. 5. 5. 6. 5. 7. 5. 5. 5. 6. 5. 6. 5. 5. 5. 5. 5. 5.
6. 6. 6. 6. 6. 6. 6. 5. 6. 6. 6. 5. 6. 6. 6. 5. 6. 5. 5. 5. 6. 5. 6. 5.
7. 5. 5. 6. 6. 5. 6. 6. 5. 5. 5. 5. 6. 5. 5. 5. 5. 6. 5. 5. 6. 5. 5. 5.
6. 6. 7. 6. 6. 6. 6. 6. 6. 6. 5. 5. 6. 5. 6. 5.]

```

Mean Absolute Error: 0.4836851598205273

Mean Squared Error: 0.3930633368592633

Root Mean Squared Error: 0.6269476348621655

```

[19]: # displaying coefficients of each feature
coefficients = pd.DataFrame(regressor.coef_,x.columns)
coefficients.columns = ['Coefficient']
print(coefficients)

```

	Coeffecient
fixed acidity	0.023274
volatile acidity	-0.991536
citric acid	-0.141268
residual sugar	0.008119
chlorides	-1.591924
free sulfur dioxide	0.005500
total sulfur dioxide	-0.003542
density	-6.069166
pH	-0.406325
sulphates	0.823603
alcohol	0.294181

[]: